

ViennaData 1.0.1 ---

User Manual



Institute for Microelectronics
Gußhausstraße 27-29 / E360
A-1040 Vienna, Austria



Copyright © 2011, Institute for Microelectronics, TU Vienna.

Main Contributors:

Karl Rupp

Current Maintainers:

Karl Rupp

Institute for Microelectronics
Vienna University of Technology
Gußhausstraße 27-29 / E360
A-1040 Vienna, Austria/Europe

Phone +43-1-58801-36001
FAX +43-1-58801-36099
Web <http://www.iue.tuwien.ac.at>

Contents

Introduction	1
1 Installation	3
1.1 Dependencies	3
1.2 Generic Installation of ViennaData	3
1.3 Building the Examples and Tutorials	3
2 Basic Usage	6
2.1 Access	6
2.2 Copy	8
2.3 Move	9
2.4 Erase	9
2.5 Query Data Availability	9
3 Advanced Usage	11
3.1 Retrieving IDs from Objects	11
3.2 Dense Data	12
3.3 Type-Based Key Dispatch	13
3.4 Default Data Types for Keys	14
4 Benchmark Results	15
4.1 ViennaData Storage Models	15
4.2 ViennaData vs. Traditional OOP	16
5 Use Case: Scientific Simulation	19
5.1 Handling Material Properties	19
5.2 Use with Discretization Methods	20
6 Library Internals	22
6.1 Main Classes	22
6.2 viennadata::all	23

7 Pitfalls	25
7.1 Object Lifetimes	25
7.2 Polymorphism	26
8 Design Decisions	27
8.1 Two Pairs of Parentheses	27
8.2 KeyType and DataType Template Parameters	28
Change Logs	30
License	31
Bibliography	32

Introduction

When designing classes in object-oriented programming languages, one has to distinguish between intrinsic properties which are inherent to the modelled entity, and data which is associated with the entity. For example, a triangle is intrinsically defined by three disjoint points, while associated data can be color or texture information in graphical applications, or material properties for two-dimensional scientific simulations.

While associated data depends on the respective application, intrinsic properties do not. Therefore, when designing reusable classes, associated data must not be part of the class itself. Reconsidering the example of a triangle, a class member for color information would immediately result in a potentially significant overhead for scientific simulations and thus reduce the reuseability of the class. However, the color information for the triangle still has to be accessible for graphical applications. This is where `ViennaData` comes into play, because it provides a convenient means to store and access associated data for small, reusable classes with intrinsic data members only.

As an example, suppose there exists a minimalistic class for triangles called `Triangle`, which provides access to its vertices. For objects of this triangle class, color information modelled by a class `RGBColor` should be stored within some application. The following lines accomplish that for a triangle object `tri` using `ViennaData`:

```
//store RGB color information:
viennadata::access<string, RGBColor>("color")(tri) = RGBColor(255,42,0);

//get color information at some other point in the application:
RGBColor col = viennadata::access<string, RGBColor>("color")(tri);
```

Data is associated with objects by keys of user-defined type in `ViennaData`. In the example above, the color information is stored using a key `"color"` of type `std::string`. However, the user is free to use any key of any type¹. Thus, any associated data object `data` of a generic type `DataType` can be stored and accessed for an object of any type using a key `key` of type `KeyType` by (cf. Chapter 2)

```
//store 'data' for obj
viennadata::access<KeyType, DataType>(key)(obj) = data;

//retrieve the stored data:
DataType data2 = viennadata::access<KeyType, DataType>(key)(obj);
```

Therefore, `ViennaData` allows to design and work with classes which focus on their intrinsic properties only, and ensures uniformity of associated data access over an arbitrarily large set of objects of different type. This is in particular important, if several different libraries with different coding styles are used within a project.

¹Some technical restrictions apply, see Sec. 2.1.

With little additional configuration (cf. Chapter 3) the code for color information access using `ViennaData` can be reduced to

```
//store RGB color information
viennadata::access<color_info>() (tri) = RGBColor(255, 42, 0);

//get color information at some other point in the application:
RGBColor col = viennadata::access<color_info>() (tri);
```

This enables uniform, type-safe data access throughout the whole application. As shown in Chapter 4, there is virtually no abstraction penalty for this gain in flexibility and uniformity. Explicit use cases are given in Chapter 5.

Internally, `ViennaData` stores data in tree-like structures, for which the object address is used as key for storing the data. However, if objects provide a means to identify themselves by numbers in the range $0, \dots, N$ for some N , the internal datastructure of `ViennaData` can be conveniently customized to a more appropriate storage scheme. Similarly, if it is a-priori known that for a particular key type only one key object per key type is used, the internal datastructure of `ViennaData` can be adapted appropriately. Chapter 6 explains the internals.

Finally, a few words are to be spent on the pitfalls and design decisions. Since data is associated with objects, object lifetimes have to be kept in mind. If an object, for which associated data is stored, reaches the end of its lifetime, the user has to care for the associated data separately. This can be very simple and fully transparent for the user in some cases, and rather tough in other cases. Moreover, `ViennaData` may lead to unexpected results when used with polymorphism. More details are given in Chapter 7, and some key design decisions are discussed in Chapter 8.

Chapter 1

Installation

This chapter shows how `ViennaData` can be integrated into a project and how the examples are built. The necessary steps are outlined for several different platforms, but we could not check every possible combination of hardware, operating system, and compiler. If you experience any trouble, please write to the mailing list at

`viennadata-support@lists.sourceforge.net`

1.1 Dependencies

- A recent C++ compiler (e.g. GCC version 4.2.x or above and Visual C++ 2008 are known to work)
- CMake [1] as build system (optional, but recommended for building the examples)

1.2 Generic Installation of ViennaData

Since `ViennaData` is a header-only library, it is sufficient to copy the `viennadata/` folder either into your project folder or to your global system include path. On Unix based systems, this is often `/usr/include/` or `/usr/local/include/`.

On Windows, the situation strongly depends on your development environment. We advise to consult the documentation of the compiler on how to set the include path correctly. With Visual Studio this is usually something like `C:\Program Files\Microsoft Visual Studio 9.0\VC\include` and can be set in Tools -> Options -> Projects and Solutions -> VC++-Directories.

1.3 Building the Examples and Tutorials

For building the examples, we suppose that CMake is properly set up on your system. The other dependencies are listed in Tab. 1.1.

File	Purpose
tutorial/access.cpp	Shows the use of <code>access</code>
tutorial/copy.cpp	Shows the use of <code>copy</code>
tutorial/custom_id.cpp	Shows how to use custom object identification
tutorial/default_data_type.cpp	Shows how a default data type for a key type is specified
tutorial/dense_data.cpp	Shows how to enable a dense storage scheme
tutorial/erase.cpp	Shows the use of <code>erase</code>
tutorial/find.cpp	Shows the use of <code>find</code>
tutorial/key_dispatch.cpp	Shows how to switch between compile time and run time key dispatch
tutorial/move.cpp	Shows the use of <code>move</code>
benchmarks/dense.cpp	Benchmarks for dense storage schemes
benchmarks/oop.cpp	Benchmarks for OOP-type class layout
benchmarks/sparse-address.cpp	Benchmarks for the default sparse storage scheme
benchmarks/sparse-custom.cpp	Benchmarks for a sparse storage scheme with custom object identification

Table 1.1: Overview of the examples in the `examples/` folder

1.3.1 Linux

To build the examples, open a terminal and change to:

```
$> cd /your-ViennaData-path/build/
```

Execute

```
$> cmake ..
```

to obtain a Makefile and type

```
$> make
```

to build the examples. If desired, one can build each example separately instead:

```
$> make access          #builds the access tutorial
$> make dense          #builds dense storage benchmark
```

Speed up the building process by using jobs, e.g. `make -j4`.



1.3.2 Mac OS X

The tools mentioned in Section 1.1 are available on Macintosh platforms too. For the GCC compiler the Xcode [2] package has to be installed. To install CMake, external portation tools such as Fink [3], DarwinPorts [4], or MacPorts [5] have to be used.

The build process of ViennaData is similar to Linux.

1.3.3 Windows

In the following the procedure is outlined for Visual Studio: Assuming that an OpenCL SDK and CMake is already installed, Visual Studio solution and project files can be created using CMake:

- Open the CMake GUI.
- Set the ViennaData base directory as source directory.
- Set the build/ directory as build directory.
- Click on 'Configure' and select the appropriate generator (e.g. Visual Studio 9 2008)
- Click on 'Generate' (you may need to click on 'Configure' one more time before you can click on 'Generate')
- The project files can now be found in the ViennaData build directory, where they can be opened and compiled with Visual Studio (provided that the include and library paths are set correctly, see Sec. 1.2).

Chapter 2

Basic Usage

In this section the fundamental operations in `ViennaData` are explained. All functionality is available out of the box and operations are checked in order to detect potential pitfalls during development. If `ViennaData` access turns out to be performance critical, several optimizations can be triggered, which is discussed in Chapter 3.

In order to use `ViennaData`, the main header file has to be included in the respective source files:

```
#include "viennadata/api.hpp"
```

No other files from `ViennaData` need to be included.

2.1 Access

The main function in `ViennaData` is the `access()` function. From a user perspective, the call

```
viennadata::access<KeyType, DataType>(key)(obj)
```

returns a reference to the data of type `DataType` stored for an arbitrary object `obj` and identified by the `key` object of type `KeyType`. At first sight, the offered functionality can be seen as a `std::map<KeyType, DataType>` for each object `obj`. However, the `access` function offers a lot more flexibility, which will be discussed in the following.

As a first example, to store RGB color information on objects `obj1`, `obj2`, and `obj3` using a `std::string`, say `"color"`, as key, the code

```
viennadata::access<string, RGBColor>("color")(obj1) = RGBColor(255,42,0);  
viennadata::access<string, RGBColor>("color")(obj2) = RGBColor(0,42,255);  
viennadata::access<string, RGBColor>("color")(obj3) = RGBColor(10,20,33);
```

is sufficient. Note that `obj1`, `obj2`, and `obj3` can be of different type. The data can be retained directly via

```
RGBColor col1 = viennadata::access<string, RGBColor>("color")(obj1);  
RGBColor col2 = viennadata::access<string, RGBColor>("color")(obj2);  
RGBColor col3 = viennadata::access<string, RGBColor>("color")(obj3);
```

In addition, several different colors can be associated with a single object `obj1` as

```
viennadata::access<string, RGBColor>("front")(obj1) = RGBColor(255,42,0);
viennadata::access<string, RGBColor>("back")(obj1) = RGBColor(12,21,255);
viennadata::access<string, RGBColor>("side")(obj1) = RGBColor(30,20,10);
```

and similarly for `obj2` and `obj3`, if desired. Since any type which fulfills the `LessThanComparable` concept ¹ (refer for example to [6]) can be used as key type, one can also use keys of different types, e.g.

```
viennadata::access<string, RGBColor>("front")(obj1) = RGBColor(255,42,0);
viennadata::access<long, RGBColor>( 12345 )(obj1) = RGBColor(12,21,255);
viennadata::access<double, RGBColor>( 42.42 )(obj1) = RGBColor(30,20,10);
```

If color is to be saved in different color spaces, this can be accomplished by adjusting the data type accordingly:

```
viennadata::access<string, RGBColor>("front")(obj1) = RGBColor(255,42,0);
viennadata::access<long, HSVColor>( 12345 )(obj1) = HSVColor(12,21,255);
viennadata::access<double, CMYKColor>(42.42 )(obj1) = CMYKColor(0,1,2,3);
```

At this point the second template argument appears to be somewhat redundant. However, please read about the design decisions in Chapter 8 in order to get more details on why the second template parameter is still required.

In the last examples, only one key object per key type has been used. As will be shown in Chapter 3, one can completely turn off run-time comparisons of keys, so that one can write

```
viennadata::access<front_color, RGBColor>()(obj1) = RGBColor(255,42,0);
viennadata::access<back_color, HSVColor>()(obj1) = HSVColor(12,21,255);
viennadata::access<side_color, CMYKColor>()(obj1) = CMYKColor(0,1,2,3);
```

where `front_color`, `back_color` and `side_color` are empty tag classes [7, 8], which are solely used as a key type. By using the types of the key only, the run time key comparisons are modified to compile time dispatches. This results in considerably faster execution, provided that it is already known at compile time, which key has to be used.

Compile time dispatches are disabled by default. On details how to enable them, refer to Sec. 3.3



With compile time dispatches enabled, the stored data can be retained in the same way:

```
RGBColor col1 = viennadata::access<front_color, RGBColor>()(obj1);
HSVColor col2 = viennadata::access<back_color, HSVColor>()(obj1);
CMYKColor col3 = viennadata::access<side_color, CMYKColor>()(obj1);
```

For the case that it is a-priori known that for each `KeyType` only one `DataType` will be used, `ViennaData` can be configured to explicitly bind the `KeyType` to a particular `DataType`. In the previous code snippet, the `front_color` can be bound to the `RGBColor` type, and similarly the `back_color` to `HSVColor` and `side_color` to `CMYKColor`. In this case, the previous code snippet reduces to

```
RGBColor col1 = viennadata::access<front_color>()(obj1);
HSVColor col2 = viennadata::access<back_color>()(obj1);
CMYKColor col3 = viennadata::access<side_color>()(obj1);
```

¹`operator<` can be used for comparisons

Similarly, no `DataType` is used for storing data in this case. However, one can still provide the `DataType` template argument as before, so existing code remains valid when a default argument is specified at some later point. More details can be found in Sec. 3.4.

2.2 Copy

In order to copy data of type `DataType` stored using a key of type `KeyType` from an object `obj_src` to an object `obj_dest`, the `copy()` function can be used as follows;

```
viennadata::copy<KeyType, DataType>(key)(obj_src, obj_dest);
```

In this case only the particular data for the particular key is copied. If the `key` argument is omitted, i.e.

```
viennadata::copy<KeyType, DataType>()(obj_src, obj_dest);
```

all data stored for `obj_src` using any key of type `KeyType` is copied to `obj_dest`. As an example, consider

```
viennadata::access<string, RGBColor>("front")(obj_src) = RGBColor(255,0,0);
viennadata::access<string, RGBColor>("back")(obj_src) = RGBColor(0,255,0);
viennadata::access<string, RGBColor>("side")(obj_src) = RGBColor(0,0,255);

//copy all data from obj_src to obj_dest:
viennadata::access<string, RGBColor>()(obj_src, obj_dest);

//access the front color (result will be RGBColor(255, 0, 0))
RGBColor red = viennadata::access<string, RGBColor>("front")(obj_dest);
```

In this example, the colors stored for `obj_src` are copied to `obj_dest` and can then be accessed using the keys `"front"`, `"back"` and `"side"`.

So far, every pair of `KeyType` and `DataType` has to be copied by a separate function call. For larger projects where multiple different keys and data types are used, this is certainly tedious. To copy all data of possibly different type for a particular key, one can use the dedicated `viennadata::all` specifier as follows:

```
viennadata::copy<KeyType, viennadata::all>()(obj_src, obj_dest);
```

Similarly, to copy all data of type `DataType` stored for `obj_src` to `obj_dest`, the `all` type can be provided for the key type as

```
viennadata::copy<viennadata::all, DataType>()(obj_src, obj_dest);
```

Finally, to copy all data of arbitrary type stored using any key of any type for `obj_src` to `obj_dest`, it is sufficient to write

```
viennadata::copy<viennadata::all, viennadata::all>()(obj_src, obj_dest);
```

`obj_src` and `obj_dest` have to be of the same type if the `viennadata::all` type is supplied.



2.3 Move

Data of type `DataType` associated with an object `obj_src` using a key of type `KeyType` can be transferred to another object `obj_dest` by using the `move()` function as follows:

```
viennadata::move<KeyType, DataType>(key)(obj_src, obj_dest);
```

The difference to the `copy()` function in the previous section is that with `move()` the data is then only available for `obj_dest`.

Similar to `copy()`, the following variants are available:

```
viennadata::move<KeyType, DataType>()(obj_src, obj_dest);
viennadata::move<KeyType, viennadata::all>()(obj_src, obj_dest);
viennadata::move<viennadata::all, DataType>()(obj_src, obj_dest);
viennadata::move<viennadata::all, viennadata::all>()(obj_src, obj_dest);
```

Since their use is similar to that of the `copy()` variants, the reader is referred to the previous section.

`obj_src` and `obj_dest` have to be of the same type if the `viennadata::all` type is supplied.



2.4 Erase

In order to delete data of type `DataType` stored using a key of type `KeyType` for an object `obj`, the `erase()` function can be used as follows:

```
viennadata::erase<KeyType, DataType>(key)(obj);
```

Similar to `copy()` and `move()`, the following variants are available:

```
viennadata::erase<KeyType, DataType>()(obj);
viennadata::erase<KeyType, viennadata::all>()(obj);
viennadata::erase<viennadata::all, DataType>()(obj);
viennadata::erase<viennadata::all, viennadata::all>()(obj);
```

They allow to delete larger amounts of data associated with `obj` by a single call. The use of `viennadata::all` is analogous to the use with `copy()` as described in Sec. 2.2.

2.5 Query Data Availability

`viennadata::find()` allows to check whether data is stored for a particular object. If there is data stored, a pointer to the data is returned. Otherwise, a `NULL` pointer is returned. The following code snippets shows how to check for data of type `DataType` associated with `obj` using a key of type `KeyType`:

```
if (viennadata::find<KeyType, DataType>(key)(obj))
    /* data available */
else
    /* data not yet available */
```

It is important to emphasize the difference to `viennadata::access`: While `find()` only checks for data availability, `access()` creates default-initialized data, if no data has been stored yet. Thus, using `viennadata::access` to query for data existence by iterating through all objects would create data of type `DataType` for every object. This overhead is avoided with `viennadata::find`, thus it is very attractive, if there is a large number of objects with only a few carrying data.

Chapter 3

Advanced Usage

While Chapter 2 outlines the basic usage of `ViennaData`, this chapter focuses on customizing the library to the user's demands and the particular environment. It is explained how the object identification can be changed, how the internal storage of data can be customized and data accesses can be tuned for maximum speed.

3.1 Retrieving IDs from Objects

In order to associate data with objects, one has to distinguish between different objects. The only portable identification mechanism is the address of an object, thus this is the default identification mechanism `ViennaData`. However, object addresses can be randomly scattered over the address space, which leads to tree-based data structures in order to provide a reasonably fast lookup of data.

However, it can be the case that there are additional means for identification provided. Let us consider a class `MyClass` equipped with a member function `id()`, which is used to identify objects. The following code configures `ViennaData` to use this member function as identification mechanism for all objects

```
namespace viennadata
{
    namespace config
    {
        template <>
        struct object_identifier<MyClass>
        {
            typedef object_provided_id    tag;
            typedef long                   id_type;

            static id_type get(MyClass const & obj) { return obj.id(); }
        };
    }
}
```

Summing up, a specialization for the template class `object_identifier` has to be provided for each class, containing the following:

- A type definition for a type `tag`. This line is identical for all user-provided overloads.

- A type definition for a type `id_type`. This specifies the type of the ID.
- A static member function `get`, which returns the ID from an object.

The specialization can be placed at any point in the source code.

3.2 Dense Data

By default, `ViennaData` internally uses a tree-based layout similar to

```
std::map<ObjType *, DataContainer>
```

to store and retrieve data for objects of type `MyClass` by the use of a data container type `DataContainer`. This induces lookup-costs of the order $\log N$, where N denotes the number of objects for which data is stored.

However, if the range of IDs of objects of a type `MyClass` is known to be in the range $[0, \dots, N]$ for some reasonable value of N , a tree-based layout can be configured to be of type

```
std::vector<DataContainer>
```

This requires that a custom identification mechanism is set up as outlined in the previous section. Moreover, the custom `id_type` must be compatible with being used as an array index. Since such a storage scheme is only memory efficient, if the number of objects is comparable to N , we refer to this scheme as *dense data storage*. To change the default tree-based data storage scheme to the dense storage scheme for objects of type `MyClass`, the lines

```
namespace viennadata
{
    namespace config
    {
        template <typename KeyType, typename DataType>
        struct storage<KeyType, DataType, MyClass>
        {
            typedef dense_data_tag    tag;
        };
    }
}
```

are sufficient. Due to the template specialization, a fine-grained control over which storage to use for the various types of data and keys is readily provided. Thus, it is possible to use both dense and sparse storage schemes simultaneously for associating data with a particular object.

For convenience, a number of macros is provided, which simplify the overloading process:

```
VIENNADATA_ENABLE_DENSE_DATA_STORAGE_FOR_KEY_DATA_OBJECT(...)
VIENNADATA_DISABLE_DENSE_DATA_STORAGE_FOR_KEY_DATA_OBJECT(...)

VIENNADATA_ENABLE_DENSE_DATA_STORAGE_FOR_KEY_DATA(...)
VIENNADATA_DISABLE_DENSE_DATA_STORAGE_FOR_KEY_DATA(...)
VIENNADATA_ENABLE_DENSE_DATA_STORAGE_FOR_KEY_OBJECT(...)
VIENNADATA_DISABLE_DENSE_DATA_STORAGE_FOR_KEY_OBJECT(...)
```



```

VIENNADATA_ENABLE_DENSE_DATA_STORAGE_FOR_DATA_OBJECT(...)
VIENNADATA_DISABLE_DENSE_DATA_STORAGE_FOR_DATA_OBJECT(...)

VIENNADATA_ENABLE_DENSE_DATA_STORAGE_FOR_KEY(...)
VIENNADATA_ENABLE_DENSE_DATA_STORAGE_FOR_DATA(...)
VIENNADATA_ENABLE_DENSE_DATA_STORAGE_FOR_OBJECT(...)

```

The parameter to be passed to the macros are the key type, the data type and/or the object type as indicated by the macro name tail. Note that the macros have to be placed in global namespace and the template overloading rules apply.

Similar to tree-based schemes, dense data storage also offers automatic memory management. However, to reduce reallocations, the user may reserve the required memory manually. To reserve data memory for data of type `DataType` with keys of type `KeyType` for up to 42 objects of the same type as `obj`, one writes

```
viennadata::reserve<KeyType,DataType>(42)(obj);
```

Similarly, memory for other data, key and object types is reserved.

3.3 Type-Based Key Dispatch

As already shown in Chapter 2, `ViennaData` can be configured to perform a purely type-based key dispatch for data access. The impact on performance can be considerable, especially when compared to potentially costly string comparisons. To enable a compile time dispatch for a key of type `SomeKeyType`, the lines

```

namespace viennadata
{
    namespace config
    {
        template <>
        struct key_dispatch<SomeKeyType>
        {
            typedef type_key_dispatch_tag    tag;
        };
    }
}

```

are sufficient. One may then omit the key parameter for the API functions provided by `ViennaData`, cf. Chapter 2.

A convenience macro is again provided:

```
VIENNADATA_ENABLE_TYPE_BASED_KEY_DISPATCH(...)
```

where the parameter is the the key type. Note that the macro has to be placed in global namespace.

When using a compile time key dispatch based on types, the key type is not required to fulfill the `LessThan-Comparable` concept.



3.4 Default Data Types for Keys

For the cases where only one data type per key type is used, `ViennaData` can be provided with a default data type for each key type. This is especially common, when using a type-based key dispatch for data access. Consider for example

```
viennadata::access<front_color, RGBColor>() (obj1) = RGBColor(255,42,0);  
viennadata::access<back_color, RGBColor>() (obj1) = RGBColor(0,42,255);  
viennadata::access<edge_color, RGBColor>() (obj1) = RGBColor(255,0,42);
```

Since the key classes are already named appropriately and always the same data type `RGBColor` is used, setting `RGBColor` as default data type for the respective keys allows to make the code more compact. This can be enabled along the lines

```
namespace viennadata  
{  
    namespace config  
    {  
        template <>  
        struct default_data_for_key<front_color>  
        {  
            typedef RGBColor    type;  
        };  
        // and similarly for back_color, edge_color  
    }  
}
```

Now, the initial code snippet can be written more compactly as

```
viennadata::access<front_color>() (obj1) = RGBColor(255,42,0);  
viennadata::access< back_color>() (obj1) = RGBColor(0,42,255);  
viennadata::access< edge_color>() (obj1) = RGBColor(255,0,42);
```

If no default data type is specified, but the second argument is nevertheless omitted, a compile time error is thrown.

A convenience macro for the default data type is provided by

```
VIENNADATA_ENABLE_DEFAULT_DATA_TYPE_FOR_KEY(...)
```

where the two parameters are the key type and the data type. Note that the macro has to be placed in global namespace.

Chapter 4

Benchmark Results

Since `ViennaData` introduces a layer for handling associated data, it is rather natural to ask for any potential runtime overhead. The aim of this chapter is to give a precise answer. Tests have been carried out on a standard desktop computer running Ubuntu 10.10.

4.1 ViennaData Storage Models

There are the following three different object-to-data models available in `ViennaData`:

- Object identification via object addresses, sparse storage.
- Custom identification via dedicated `id()` member, sparse storage.
- Custom identification via dedicated `id()` member, dense storage (cf. Section 3.2).

In addition, the following data key dispatch mechanisms are compared in the following:

- Key dispatch at run time for keys of type `std::string`.
- Key dispatch at run time for keys of type `long`.
- Key dispatch at compile time (cf. Section 3.3).

This results in nine different configurations for the subsequent tests. All comparisons are carried out for objects of the following class:

```
class SlimClass
{
public:
    SlimClass(double v = 1.0, size_t i = 0) : value_(v), id_(i) {}

    double value() const { return value_; }
    size_t id() const { return id_; }
private:
    double value_;
    size_t id_;
};
```

	Key Dispatch		
Identification	std::string	long	compile time
&obj, sparse	84	53	43
obj.id(), sparse	83	52	43
obj.id(), dense	48	15	4

Table 4.1: Execution time for summing up data of 1 000 objects of type `SlimClass`, repeated 1 000 times. Execution times in milliseconds.

	Key Dispatch		
Identification	std::string	long	compile time
&obj, sparse	333	247	173
obj.id(), sparse	318	228	178
obj.id(), dense	151	73	5

Table 4.2: Execution time for summing up data of 1 000 000 objects of type `SlimClass`. Execution times in milliseconds.

The first test with execution times given in Table 4.1 compares the execution times of summing up random values stored and retrieved for 1 000 objects of type `SlimClass` using `viennadata::access<>()`, repeated 1 000 times.

Let us first consider the differences among the three rows. As can clearly be seen, a change of the object identification from address-based to using the member function `id()` does not change execution speed, if the internal storage scheme remains tree-based (*sparse*). However, if the `id()` member is used for a dense storage scheme as described in Section 3.2, execution times are reduced by around 40 milliseconds irrespective of the key dispatch method employed.

Comparing the different columns in Table 4.1, it can be seen that a string-based dispatch is most expensive as expected. A dispatch based on integers is faster, but the differences only become significant, if a dense storage scheme is used. The compile-time dispatch is notably faster than dispatches at run time. In particular, execution times differ by a factor of four and twelve for the dense storage scheme respectively.

In a second benchmark, the same procedure is carried out for 1 000 000 objects without repetition. Thus, a naive guess is that execution times should be roughly the same. However, as Table 4.2 shows, the logarithmic lookup times in a sparse, tree-based storage layout leads to much higher execution times than for the first test. This may be also partly due to less efficient CPU cache use caused by the larger amount of data involved.

Finally, it should be emphasized that in real applications there is typically much less emphasis on data retrieval as compared to our benchmarks. Thus, even though a sparse storage scheme is slower than a dense storage scheme, it may be fast enough for most applications.

4.2 ViennaData vs. Traditional OOP

As discussed in the introduction already, `ViennaData` allows to design small, lightweight classes with high reusability. In this section we compare data access times of `ViennaData`

Setup	Exec. Time (us), 10^3 Objects	Exec. Time (ms), 10^6 Objects
ViennaData, config 1	84	333
ViennaData, config 2	52	228
ViennaData, config 3	4	5
OOP, FatClass<10>	1.3	4
OOP, FatClass<100>	2.1	11
OOP, FatClass<1000>	2.5	11

Table 4.3: Execution time for summing up one `double` value from each object.

with other classes that stem from traditional OOP design and thus have a higher number of class member variables.

The objects for which data is attached using `ViennaData` are of type `SlimClass` as introduced in the previous section. The following three access models are used:

1. Object identification via object addresses, sparse storage, key dispatch at run time for keys of type `std::string`.
2. Custom identification via dedicated `id()` member, sparse storage, key dispatch at run time for keys of type `long`.
3. Custom identification via dedicated `id()` member, dense storage, key dispatch at compile time.

Note that the three models represents the diagonals in Table 4.1 and Table 4.2.

For the traditional OOP classes, we use the following template class

```
template <size_t num_bytes>
class FatClass
{
public:
    FatClass(double v = 1.0, size_t i = 0) : value_(v), id_(i) {}

    double value() const { return value_; }
    size_t id() const { return id_; }
private:
    double value_;
    char payload[num_bytes];
    size_t id_;
};
```

where the template parameter `num_bytes` denotes the number of bytes used for other data members. `FatClass<10>`, `FatClass<100>` and `FatClass<1000>` are used.

As can be seen in Table 4.3, access times with dense storage are by a factor of two larger for 1 000 objects, but by a factor of two smaller when using 1 000 000 objects. Thus, `ViennaData` allows to achieve even better execution performance than traditional OOP for a large number of objects. The reason is that the relevant data is located in a contiguous piece of memory when using `ViennaData` with a dense data layout, thus offering better caching possibilities. On the other hand, with standard OOP the memory layout is such that relevant data for each object is followed by payload data and this pattern is repeated for each

object. Thus, relevant data is scattered over a larger memory region, reducing caching possibilities.

The sparse storage layout is not competitive in terms of access times as the number of objects grows. However, data access times are often negligible compared to data manipulation, thus the increased data access time may be in many cases still small on the overall.

Chapter 5

Use Case: Scientific Simulation

The potential of `ViennaData` for scientific simulations is discussed in this chapter. More precisely, the use of `ViennaData` for the solution of partial differential equations by means of discretization methods such as the Finite Element Method or the Finite Volume Method is presented.

It is assumed in the following that the simulation domain, which can be one-, two- or three-dimensional, is decomposed into small elements, say lines, triangles or tetrahedrons. The simulation domains can represent arbitrary objects, for instance buildings, human bones or microelectronic devices, and the physical effects simulated can be displacements, stresses, heat, charge densities, etc.

5.1 Handling Material Properties

Let us consider the classes `Line`, `Triangle`, and `Tetrahedron`, which reflect the respective geometric primitives. When running a two-dimensional simulation, material properties have to be associated with triangles, thus one might be tempted to add a `material` member variable of some type, say, `Material` to the `Triangle` class. However, reusing the same class for facets of a `Tetrahedron` in three-dimensional simulations would result in an unused member variable `material` in `Triangle`, because materials are associated with a `Tetrahedron` in three-dimensional simulations. Similarly, material information should not be a member of a `line`.

In order to preserve reusability of the `Line`, `Triangle` and `Tetrahedron` classes, `ViennaData` can be used as follows for the three-dimensional case:

```
Tetrahedron tet;
/* initialize tet */

//store material on tet using ViennaData:
viennadata::access<string, Material>("material")(tet) = Material("wood");
```

Then, at any later point, the material for `tet` can be conveniently retrieved by

```
Material tet_mat = viennadata::access<string, Material>("material")(tet);
```

Other keys can be used as described in Chapter 2. Note that one will typically have thousands of cells, not just a single one.

In order to allow a uniform material retrieval independent of the spatial dimension used for the simulation, one can provide a **typedef** for the respective cell type:

```
typedef Tetrahedron      CellType; //for 3d
//typedef Triangle      CellType; //for 2d
//typedef Line           CellType; //for 1d

CellType cell;

/* Initialize cell here */

//store data on cell:
viennadata::access<string, Material>("material")(cell) = Material("wood");

//retrieve material from cell at some later point:
Material c_mat = viennadata::access<string, Material>("material")(cell);
```

Thus, in this setting `ViennaData` provides unified access for objects of type `Line`, `Triangle` and `Tetrahedron`. Moreover, one is not restricted to a particular implementation and can still exchange the implementations for the geometric primitives at some later point without affecting the storage of materials.

By using `ViennaData` one can easily exchange the implementation of the objects the data is associated with.



5.2 Use with Discretization Methods

The various discretization schemes such as the Finite Volume Method or the Finite Element Method associate unknowns with vertices, lines, facets, and/or cells within a mesh. A similar reasoning as for the storage of material properties applies to the storage of unknowns (or their identification indices) as well. Moreover, one may not simulate a single effect, but several, which may or may not be coupled. For instance, one may want to simulate mechanical stresses as well as heat flow in a building. Thus, data required for the different simulations may have to be separated. Such a scenario would be almost impossible to handle, if the unknown identification indices were members of the respective classes `Line`, `Triangle` and `Tetrahedron`. However, with `ViennaData` unknown indices can be stored and retrieved for a point object `vertex` as

```
//store unknown index id_1 for simulation 1 (long sim_1 = 1;):
viennadata::access<long, UnknownIndex>(sim_1)(vertex) = id_1;
//store unknown index id_2 for simulation 2 (long sim_2 = 2;):
viennadata::access<long, UnknownIndex>(sim_2)(vertex) = id_2;

/* lots of other computations done here */

//during simulation 1:
UnknownIndex id_1 = viennadata::access<long, UnknownIndex>(sim_1)(vertex);
//during simulation 2:
UnknownIndex id_2 = viennadata::access<long, UnknownIndex>(sim_2)(vertex);
```

Thus, not only remains the implementation of geometric primitives light-weight and reusable, the resulting code due to `ViennaData` is very compact and expressive. To obtain high exe-

cution performance, a dense storage scheme should be employed for accessing the unknown indices, cf. Section 3.2.

Another important aspect for scientific simulation is the handling of boundary conditions. Typically only a subset of elements on the boundary is effected by a certain type of boundary conditions, thus it would be highly inefficient to use a boolean member variable for all elements in the mesh. Here, the default sparse storage scheme of `ViennaData` turns out to be memory-efficient, when for instance flagging boundary triangles:

```
// tag tri1, tri42 and tri360 for boundary conditions in simulation 1:
viennadata::access<long, bool>(sim_1)(tri1)   = true;
viennadata::access<long, bool>(sim_1)(tri42)  = true;
viennadata::access<long, bool>(sim_1)(tri360) = true;
// tag only tri1 and tri360 for boundary conditions in simulation 1:
viennadata::access<long, bool>(sim_2)(tri1)   = true;
viennadata::access<long, bool>(sim_2)(tri360) = true;

/* lots of other computations done here */

// iterate over triangles using tri_iter in simulation 1
// and check for boundary conditions:
if ( viennadata::find<long, bool>(sim_1)(*tri_iter) )
    //boundary triangle for simulation 1
else
    //no boundary

// iterate over triangles using tri_iter in simulation 1
// and check for boundary conditions:
if ( viennadata::find<long, bool>(sim_2)(*tri_iter) )
    //boundary triangle for simulation 2
else
    //no boundary
```

Note that by using `viennadata::access` in the if-clauses, a boolean would be default-constructed and stored on each triangle. By using `viennadata::find`, only three and two booleans are stored for the two simulations respectively.

`ViennaData` provides a flexible, unified interface for both sparse and dense storage schemes.



Chapter 6

Library Internals

In order to configure `ViennaData` in the best way for the respective project and to avoid common pitfalls, it is of advantage to understand some of the library-internals. The focus is on explaining the *big picture* rather than in going through all the details. These details can either be found in the Doxygen-generated documentation, or directly in the source code.

6.1 Main Classes

All data resides in the instances of the template class

```
template <typename KeyType,  
          typename DataType,  
          typename ObjectType>  
class data_container;
```

located in `data_container.hpp`. As can be seen from the list of template parameters, each triple of `KeyType`, `DataType`, and `ObjectType` lets the compiler create another instance of `data_container`. Thus, compilation times will increase with the number of different types involved.

`data_container` uses a singleton pattern and is directly interfaced by the interface functions

```
viennadata::access<KeyType,  DataType> () ()  
viennadata::copy<KeyType,   DataType> () ()  
viennadata::move<KeyType,   DataType> () ()  
viennadata::erase<KeyType,  DataType> () ()  
viennadata::find<KeyType,   DataType> () ()
```

The two parentheses arguments are due to the following:

1. The inner parentheses belong to a function returning a proxy class
2. The outer parentheses denote the parenthesis operator `operator ()` of the proxy class

Within the parenthesis operator, the appropriate functions of `data_container` are called.

The storage type in `data_container` is obtained from a traits class `container` in `traits/data_container.hpp`, which provides the correct storage scheme. By default, the storage is equivalent to

```
std::map<ObjectType *, std::map<KeyType, DataType> >
```

Thus, the first call

```
viennadata::access<KeyType, DataType>(key)(obj) = data;
```

for obj of type Object is equivalent to

```
std::map<ObjectType *, std::map<KeyType, DataType> > storage;  
storage[&obj][key] = data;
```

and subsequent calls reuse the storage object. If a dense storage scheme is used (cf. Section 3.2 and config/storage.hpp), the storage type is equivalent to

```
std::vector< std::map<KeyType, DataType> >
```

and the first call to viennadata::access can be seen as

```
std::vector< std::map<KeyType, DataType> > storage;  
storage.resize(some_size);  
storage[obj.id()][key] = data;
```

The configuration of the ID retrieval (here: obj.id()) is discussed in Section 3.1 and implemented in object_identifier.hpp, while the need to call viennadata::reserve in order to provide a value for some_size is discussed in Section 3.2.

Finally, a compile time key dispatch (cf. Section 3.3) eliminates the inner map of the storage scheme, hence the data is organized as

```
std::vector< DataType > dense_storage;
```

for dense storage, and as

```
std::map<ObjectType *, DataType > sparse_storage;
```

for sparse storage. The traits classes in traits/data_container.hpp ensure that the correct access mechanism is used for the various operations.

6.2 viennadata::all

In order to make the following code lines

```
viennadata::access<std::string, long>("foo")(obj1) = 42;  
viennadata::access<long, double>(360)(obj1) = 3.1415;  
viennadata::copy<viennadata::all, viennadata::all>()(obj1, obj2);
```

work as expected, smart coding has to be applied. As outlined in the previous section, the first two lines create two storage objects

```
std::map<ObjectType *, std::map<std::string, long> >  
std::map<ObjectType *, std::map<long, double> >
```

for which the data must be transferred. However, there is no means to iterate over compiler generated types, therefore the generated types have to be tracked at runtime.

A technique called *type erasure* [7] is used to store pairs of different (!) `KeyTypes` and `DataTypes` in an array, see `key_value_pair.hpp` for details on the type erasure and the file `key_value_registration.hpp` for the array. When using the `viennadata::all` type for copy, move, or erase operations, the type-erased `KeyTypes` and `DataTypes` are iterated and compared. If the types match, a virtual function call unwraps the types and calls the appropriate `data_container`-member. This virtual function call is the reason, why the source and the destination object have to be of the same type.

Chapter 7

Pitfalls

The increased reusability of classes due to a separation of intrinsic and associated data (cf. [Introduction](#)) introduces a few pitfalls one should be aware of. Depending on the code design and the application, these pitfalls can be somewhere between completely unrealistic and painful. In the following, we discuss the most important of them.

7.1 Object Lifetimes

Consider the following code, where the class `A` is defined somewhere else:

```
A * myA = new A();  
viennadata::access<long, double>(42) (*myA) = 3.1415;  
delete myA;
```

The problem here is that data associated with `A` is not erased prior to the deletion of `myA`. Thus, there is still the double value 3.1415 stored using the address of `myA`, but there is no direct way to access that data. To avoid confusion: The data is not lost in a sense of memory leak, because the tree managing the data is still consistent and properly free'd at program exit. However, the user has no direct way to retain an object which has the same address as the one pointed to by `myA`.

In order to ensure that all data is properly deleted, when the respective object is destroyed, the following options are available:

- Manually call `viennadata::erase<KeyType, DataType>() (*myA)` for all pairs of `KeyType` and `DataType` used.
- Manually ensure that `viennadata::erase<all, all>() (*myA)` is called prior to deleting `*myA`. This requires that all data is registered properly, which is by default the case.
- Use `viennadata::free(*myA)` instead of `delete`. This will call `viennadata::erase<all, all>() (*myA)` and then `delete myA`. It will not work, if `myA` points to an array, for which `delete[]` has been called.
- If the sources of `myA` can be modified, then one can inject the call to `viennadata::erase<all, all>() (*myA)` into the destructor as follows:

```
A::~~A()
{
    viennadata::erase<all,all>()(*this);
}
```

This will automatically remove all data stored for the object, unless automatic registration of key and data types is turned off.

- Provide another identification mechanism that is not based on object lifetimes, see Section 3.1. Thus, if an object with the same ID is created at some later time, the data can be accessed again.

7.2 Polymorphism

Since `ViennaData` relies on a static type system, it cannot reasonably deal with run time polymorphism. Consider

```
class Base { };
class Derived : public Base {};

Derived d;
Base * pb = &d;
viennadata<long, double>(360)(*pb) = 3.1415;
std::cout << viennadata<long, double>(360)(d) << std::endl;
```

Even though data is stored for the same object, the resulting program will still print a zero value. The reason is that data is stored for an object of type `Base`, while it is retrieved for an object of type `Derived`.

Better support for run time polymorphism might be provided in future versions of `ViennaData`. In the current version, library users are advised to use custom identification mechanisms and provide objects of the base type only.

The use of `ViennaData` with run time polymorphism may lead to unexpected results if no custom identification mechanism is provided!



Chapter 8

Design Decisions

Every now and then a programmer is faced with the question “Why have you implemented it that way?”. This is then the chance to outline the various pros and cons of different approaches and assign appropriate weights such that the approach taken turns out to be the best.

The first question is sometimes followed or replaced by “Why haven’t you done it like this:?”. In this case, the various pros and cons of different approaches should be outlined again and the taken weights have to be justified - even if it often tends to become a rather emotional debate instead.

In this chapter the various design considerations are presented and the reasons why the library is the way it is are given. Remaining questions or inputs for further improvements should be sent to the mailing list at

`viennadata-support@lists.sourceforge.net`

8.1 Two Pairs of Parentheses

Probably the most immediate question relates to the two pairs of parentheses as in

```
viennadata::access<long, double>(42)(obj) = 360.0;
```

One could have just used a single pair as follows:

```
viennadata::access<long, double>(42, obj) = 360.0;
```

However, consider the alternatives for other functions of ViennaData:

```
//reserve space for 100 objects, keys of type long:
viennadata::reserve<long, DataType>(100, obj);

//copy data stored at key '42' for obj1 to obj2:
viennadata::copy<long, DataType>(42, obj1, obj2);

//copy all data with keys of type long:
viennadata::copy<long, DataType>(obj1, obj2);

//erase all data with keys of type long for obj:
viennadata::erase<long, DataType>(obj);
```

Even though only three different functions were used, there are four different meanings of the first parameters: First, the parameter denotes the number of objects, then it denotes the data, after that the source object, and finally the object for which data is to be erased. In particular the last line is not intuitive - it suggests that `obj` is erased.

Compare the previous snippet with the actual code:

```
viennadata::reserve<long, DataType>(100)(obj);
viennadata::copy<long,   DataType>(42)(obj1, obj2);
viennadata::copy<long,   DataType>()(obj1, obj2);
viennadata::erase<long,  DataType>()(obj);
```

Here, the first pair of parentheses always refers to the set of data manipulated, while the parameters in the second pair always refer to the objects the data refers to. In the library author's point of view, the second version is more intuitive and has thus been selected.

The first pair of parentheses always refers to the data being manipulated (either key or number of objects to be reserved), while the second pair of parentheses always refers to the objects the data is associated with.



8.2 KeyType and DataType Template Parameters

Considering the code

```
//store '360.0' for 'obj' using key 42
viennadata::access<long, double>(42)(obj) = 360.0;
```

one is tempted to ask, whether the template parameters are redundant. In this example, this is certainly the case. However, a slight variation of the code above looks like this:

```
//store '360.0' for 'obj' using key 42
viennadata::access<std::string, double>("42")(obj) = 360.0;
```

Clearly, as soon as character strings are used as key, type deduction fails. Moreover, when using compile time key dispatch, one has to supply the key type anyway.

As for the data type, which could be automatically deduced to be of type `double`, there are two problems involved: The one is that for accessing data with possibly implicit conversion, there is no other option but to supply the data type:

```
//store '360.0' for 'obj' using key 42
int value = viennadata::access<long, double>(42)(obj);
```

The other problem is due to unwanted side-effects, when changing data types while refactoring code. If the `DataType` were optional, the code

```
long data = 42;
/* maybe some other code here */

//store data for obj:
viennadata::access<std::string>("data")(obj) = data;

//at some other place, maybe even a different source file:
long data2 = viennadata::access<std::string, long>("data")(obj);
```


were legal. However, when changing the type of `data` from `long` to, say, `int`, the data retrieval would be broken. Thus, to unify the interface for access, copy, move, erase, and find, the key type and the data type always have to be supplied. The only possible exception has to be explicitly enabled by the user, cf. Section 3.4.

The key type and the data type always has to be specified, when calling functions from `ViennaData`. The only exception has to be explicitly enabled by the user.



Change Logs

Version 1.0.1

Added convenience macros for changing the storage type, for default data types and for type-based key dispatch. Moreover, extended `find()` for a key-type dispatch and fixed a compilation issue when using data of type `bool` for dense storage.

Version 1.0.0

First release

License

Copyright (c) 2011, Institute for Microelectronics, TU Wien

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

Bibliography

- [1] “CMake.” [Online]. Available: <http://www.cmake.org/>
- [2] “Xcode Developer Tools.” [Online]. Available: <http://developer.apple.com/technologies/tools/xcode.html>
- [3] “Fink.” [Online]. Available: <http://www.finkproject.org/>
- [4] “DarwinPorts.” [Online]. Available: <http://darwinports.com/>
- [5] “MacPorts.” [Online]. Available: <http://www.macports.org/>
- [6] M. H. Austern, *Generic Programming and the STL: Using and Extending the C++ Standard Template Library*. Boston, MA, USA: Addison-Wesley Longman Publishing Co., Inc., 1998.
- [7] A. Alexandrescu, *Modern C++ Design: Generic Programming and Design Patterns Applied*. Boston, MA, USA: Addison-Wesley Longman Publishing Co., Inc., 2001.
- [8] D. Vandevoorde and N. M. Josuttis, *C++ Templates*. Boston, MA, USA: Addison-Wesley Longman Publishing Co., Inc., 2002.